

# Image\_ClassificationCIFAR100

April 13, 2020

```
[0]: from __future__ import print_function
import keras
from keras.datasets import cifar100
from keras.preprocessing.image import ImageDataGenerator
from keras.models import Sequential
from keras.layers import Dense, Dropout, Activation, Flatten
from keras.layers import Conv2D, MaxPooling2D
import keras.backend as K
import tensorflow as tf

import os
import pickle
import numpy as np
import h5py
import pandas as pd
import matplotlib.pyplot as plt
```

## 0.1 Downloading DataSet

We use the dataset CIFAR100 for Multiclass Classification

- Batch\_size: amount of samples that will be feed-forward in our model at once
- Img\_width = 32, Img\_height = 32, channels = 3 as RGB image
- loss\_function: compare predictions with ground truth during training
- number of classes: 100
- optimizer: method by which we update the weights of our neural network

```
[4]: num_classes = 100
save_dir = os.path.join(os.getcwd(), 'saved_models')
model_name = 'cifar100.h5'

# The data, shuffled and split between train and test sets:
(x_train, y_train), (x_test, y_test) = cifar100.load_data()
print('x_train shape:', x_train.shape)
print(x_train.shape[0], 'train samples')
```

```

print(x_test.shape[0], 'test samples')

# Convert class vectors to binary class matrices.
y_train = keras.utils.to_categorical(y_train, num_classes)
y_test = keras.utils.to_categorical(y_test, num_classes)

x_train = x_train.astype('float32')
x_test = x_test.astype('float32')
x_train /= 255.
x_test /= 255.

```

x\_train shape: (50000, 32, 32, 3)  
50000 train samples  
10000 test samples

## 0.2 Network

```

[0]: model = Sequential()

model.add(Conv2D(128, (3, 3), padding='same',
                input_shape=x_train.shape[1:]))
model.add(Activation('elu'))
model.add(Conv2D(128, (3, 3)))
model.add(Activation('elu'))
model.add(MaxPooling2D(pool_size=(2, 2)))

model.add(Conv2D(256, (3, 3), padding='same'))
model.add(Activation('elu'))
model.add(Conv2D(256, (3, 3)))
model.add(Activation('elu'))
model.add(MaxPooling2D(pool_size=(2, 2)))
model.add(Dropout(0.25))

model.add(Conv2D(512, (3, 3), padding='same'))
model.add(Activation('elu'))
model.add(Conv2D(512, (3, 3)))
model.add(Activation('elu'))
model.add(MaxPooling2D(pool_size=(2, 2)))
model.add(Dropout(0.25))

model.add(Flatten())
model.add(Dense(1024))
model.add(Activation('elu'))
model.add(Dropout(0.5))
model.add(Dense(num_classes))
model.add(Activation('softmax'))

```

```
# initiate RMSprop optimizer
opt = keras.optimizers.rmsprop(lr=0.0001, decay=1e-6)

# training the model using RMSprop
model.compile(loss='categorical_crossentropy',
              optimizer=opt,
              metrics=['accuracy'])
```

## 1 Model Fitting

```
[8]: epochs = 50
      num_predictions = 20
      batch_size = 64
      validations = []
      history = model.fit(x_train,
                          y_train, batch_size=batch_size, epochs=epochs, validation_data=(x_test,
                          y_test), shuffle=True)
      pickle.dump(validations, open("loss_validation.p", 'wb'))
```

Train on 50000 samples, validate on 10000 samples

```
Epoch 1/50
50000/50000 [=====] - 63s 1ms/step - loss: 2.1612 -
accuracy: 0.4467 - val_loss: 2.1108 - val_accuracy: 0.4639
Epoch 2/50
50000/50000 [=====] - 63s 1ms/step - loss: 2.0398 -
accuracy: 0.4723 - val_loss: 2.0401 - val_accuracy: 0.4722
Epoch 3/50
50000/50000 [=====] - 63s 1ms/step - loss: 1.9221 -
accuracy: 0.4950 - val_loss: 1.9920 - val_accuracy: 0.4933
Epoch 4/50
50000/50000 [=====] - 63s 1ms/step - loss: 1.8114 -
accuracy: 0.5202 - val_loss: 1.9699 - val_accuracy: 0.4927
Epoch 5/50
50000/50000 [=====] - 62s 1ms/step - loss: 1.6899 -
accuracy: 0.5494 - val_loss: 1.9667 - val_accuracy: 0.4923
Epoch 6/50
50000/50000 [=====] - 62s 1ms/step - loss: 1.5766 -
accuracy: 0.5765 - val_loss: 1.9137 - val_accuracy: 0.5073
Epoch 7/50
50000/50000 [=====] - 62s 1ms/step - loss: 1.4688 -
accuracy: 0.6005 - val_loss: 1.9536 - val_accuracy: 0.5080
Epoch 8/50
50000/50000 [=====] - 62s 1ms/step - loss: 1.3625 -
accuracy: 0.6215 - val_loss: 1.8924 - val_accuracy: 0.5223
Epoch 9/50
50000/50000 [=====] - 62s 1ms/step - loss: 1.2707 -
```

accuracy: 0.6465 - val\_loss: 1.8472 - val\_accuracy: 0.5310  
 Epoch 10/50  
 50000/50000 [=====] - 62s 1ms/step - loss: 1.1744 -  
 accuracy: 0.6705 - val\_loss: 1.8820 - val\_accuracy: 0.5302  
 Epoch 11/50  
 50000/50000 [=====] - 62s 1ms/step - loss: 1.0815 -  
 accuracy: 0.6911 - val\_loss: 1.8863 - val\_accuracy: 0.5297  
 Epoch 12/50  
 50000/50000 [=====] - 62s 1ms/step - loss: 0.9893 -  
 accuracy: 0.7154 - val\_loss: 1.9413 - val\_accuracy: 0.5269  
 Epoch 13/50  
 50000/50000 [=====] - 62s 1ms/step - loss: 0.9133 -  
 accuracy: 0.7347 - val\_loss: 1.9460 - val\_accuracy: 0.5364  
 Epoch 14/50  
 50000/50000 [=====] - 63s 1ms/step - loss: 0.8345 -  
 accuracy: 0.7546 - val\_loss: 1.9120 - val\_accuracy: 0.5368  
 Epoch 15/50  
 50000/50000 [=====] - 63s 1ms/step - loss: 0.7737 -  
 accuracy: 0.7707 - val\_loss: 1.9976 - val\_accuracy: 0.5280  
 Epoch 16/50  
 50000/50000 [=====] - 63s 1ms/step - loss: 0.7010 -  
 accuracy: 0.7888 - val\_loss: 1.9605 - val\_accuracy: 0.5407  
 Epoch 17/50  
 50000/50000 [=====] - 63s 1ms/step - loss: 0.6459 -  
 accuracy: 0.8051 - val\_loss: 1.9510 - val\_accuracy: 0.5394  
 Epoch 18/50  
 50000/50000 [=====] - 63s 1ms/step - loss: 0.5878 -  
 accuracy: 0.8205 - val\_loss: 2.0290 - val\_accuracy: 0.5434  
 Epoch 19/50  
 50000/50000 [=====] - 63s 1ms/step - loss: 0.5398 -  
 accuracy: 0.8350 - val\_loss: 2.0400 - val\_accuracy: 0.5385  
 Epoch 20/50  
 50000/50000 [=====] - 63s 1ms/step - loss: 0.4978 -  
 accuracy: 0.8461 - val\_loss: 2.1446 - val\_accuracy: 0.5496  
 Epoch 21/50  
 50000/50000 [=====] - 63s 1ms/step - loss: 0.4564 -  
 accuracy: 0.8581 - val\_loss: 2.0576 - val\_accuracy: 0.5379  
 Epoch 22/50  
 50000/50000 [=====] - 63s 1ms/step - loss: 0.4289 -  
 accuracy: 0.8669 - val\_loss: 2.2093 - val\_accuracy: 0.5501  
 Epoch 23/50  
 50000/50000 [=====] - 63s 1ms/step - loss: 0.3993 -  
 accuracy: 0.8765 - val\_loss: 2.2060 - val\_accuracy: 0.5466  
 Epoch 24/50  
 50000/50000 [=====] - 63s 1ms/step - loss: 0.3683 -  
 accuracy: 0.8849 - val\_loss: 2.1630 - val\_accuracy: 0.5380  
 Epoch 25/50  
 50000/50000 [=====] - 64s 1ms/step - loss: 0.3493 -

accuracy: 0.8924 - val\_loss: 2.1684 - val\_accuracy: 0.5444  
 Epoch 26/50  
 50000/50000 [=====] - 63s 1ms/step - loss: 0.3204 -  
 accuracy: 0.9000 - val\_loss: 2.2401 - val\_accuracy: 0.5426  
 Epoch 27/50  
 50000/50000 [=====] - 63s 1ms/step - loss: 0.3044 -  
 accuracy: 0.9057 - val\_loss: 2.2343 - val\_accuracy: 0.5401  
 Epoch 28/50  
 50000/50000 [=====] - 63s 1ms/step - loss: 0.2868 -  
 accuracy: 0.9104 - val\_loss: 2.3049 - val\_accuracy: 0.5437  
 Epoch 29/50  
 50000/50000 [=====] - 63s 1ms/step - loss: 0.2707 -  
 accuracy: 0.9157 - val\_loss: 2.3941 - val\_accuracy: 0.5406  
 Epoch 30/50  
 50000/50000 [=====] - 63s 1ms/step - loss: 0.2590 -  
 accuracy: 0.9206 - val\_loss: 2.4278 - val\_accuracy: 0.5456  
 Epoch 31/50  
 50000/50000 [=====] - 63s 1ms/step - loss: 0.2407 -  
 accuracy: 0.9243 - val\_loss: 2.3932 - val\_accuracy: 0.5455  
 Epoch 32/50  
 50000/50000 [=====] - 63s 1ms/step - loss: 0.2358 -  
 accuracy: 0.9268 - val\_loss: 2.4607 - val\_accuracy: 0.5510  
 Epoch 33/50  
 50000/50000 [=====] - 63s 1ms/step - loss: 0.2301 -  
 accuracy: 0.9288 - val\_loss: 2.4346 - val\_accuracy: 0.5411  
 Epoch 34/50  
 50000/50000 [=====] - 63s 1ms/step - loss: 0.2206 -  
 accuracy: 0.9312 - val\_loss: 2.4728 - val\_accuracy: 0.5487  
 Epoch 35/50  
 50000/50000 [=====] - 63s 1ms/step - loss: 0.2079 -  
 accuracy: 0.9358 - val\_loss: 2.6035 - val\_accuracy: 0.5439  
 Epoch 36/50  
 50000/50000 [=====] - 63s 1ms/step - loss: 0.2057 -  
 accuracy: 0.9358 - val\_loss: 2.2766 - val\_accuracy: 0.5425  
 Epoch 37/50  
 50000/50000 [=====] - 63s 1ms/step - loss: 0.1968 -  
 accuracy: 0.9399 - val\_loss: 2.3900 - val\_accuracy: 0.5441  
 Epoch 38/50  
 50000/50000 [=====] - 63s 1ms/step - loss: 0.1909 -  
 accuracy: 0.9410 - val\_loss: 2.4747 - val\_accuracy: 0.5465  
 Epoch 39/50  
 50000/50000 [=====] - 64s 1ms/step - loss: 0.1879 -  
 accuracy: 0.9418 - val\_loss: 2.4327 - val\_accuracy: 0.5454  
 Epoch 40/50  
 50000/50000 [=====] - 64s 1ms/step - loss: 0.1770 -  
 accuracy: 0.9453 - val\_loss: 2.4911 - val\_accuracy: 0.5500  
 Epoch 41/50  
 50000/50000 [=====] - 63s 1ms/step - loss: 0.1757 -

```

accuracy: 0.9454 - val_loss: 2.2837 - val_accuracy: 0.5423
Epoch 42/50
50000/50000 [=====] - 63s 1ms/step - loss: 0.1706 -
accuracy: 0.9460 - val_loss: 2.5532 - val_accuracy: 0.5392
Epoch 43/50
50000/50000 [=====] - 63s 1ms/step - loss: 0.1644 -
accuracy: 0.9483 - val_loss: 2.3649 - val_accuracy: 0.5448
Epoch 44/50
50000/50000 [=====] - 63s 1ms/step - loss: 0.1651 -
accuracy: 0.9497 - val_loss: 2.3895 - val_accuracy: 0.5465
Epoch 45/50
50000/50000 [=====] - 63s 1ms/step - loss: 0.1651 -
accuracy: 0.9504 - val_loss: 2.8186 - val_accuracy: 0.5517
Epoch 46/50
50000/50000 [=====] - 64s 1ms/step - loss: 0.1588 -
accuracy: 0.9517 - val_loss: 2.8806 - val_accuracy: 0.5456
Epoch 47/50
50000/50000 [=====] - 63s 1ms/step - loss: 0.1553 -
accuracy: 0.9528 - val_loss: 2.6195 - val_accuracy: 0.5504
Epoch 48/50
50000/50000 [=====] - 64s 1ms/step - loss: 0.1518 -
accuracy: 0.9523 - val_loss: 2.6581 - val_accuracy: 0.5434
Epoch 49/50
50000/50000 [=====] - 63s 1ms/step - loss: 0.1514 -
accuracy: 0.9539 - val_loss: 2.2576 - val_accuracy: 0.5336
Epoch 50/50
50000/50000 [=====] - 63s 1ms/step - loss: 0.1499 -
accuracy: 0.9551 - val_loss: 2.1556 - val_accuracy: 0.5275

```

## 2 Plots for Categorical Cross Entropy Loss

Also called Softmax Loss. It is a Softmax activation plus a Cross-Entropy loss. If we use this loss, we will train a CNN to output a probability over the C classes for each image. It is used for multi-class classification.

```

[9]: # model.evaluate(x=x_test, y=y_test)

# Generate generalization metrics
score = model.evaluate(x=x_test, y=y_test, verbose=1)
print(f'Test loss: {score[0]} / Test accuracy: {score[1]}')

# Visualize history
# Plot history: Loss
plt.plot(history.history['val_loss'])
plt.title('Validation loss history')
plt.ylabel('Loss value')
plt.xlabel('No. epoch')


```


```
plt.show()

# Plot history: Accuracy
plt.plot(history.history['val_accuracy'])
plt.title('Validation accuracy history')
plt.ylabel('Accuracy value (%)')
plt.xlabel('No. epoch')
plt.show()
```

10000/10000 [=====] - 5s 523us/step

Test loss: 2.1555932670593263 / Test accuracy: 0.5274999737739563





```
[0]: ans=model.predict(x_test)
```

```
[11]: print(ans.shape)
```

```
(10000, 100)
```

```
[0]: final_label=np.zeros(10000)
gt=np.zeros(10000)

for i in range(10000):
    final_label[i]=np.argmax(ans[i,:])
    gt[i]=np.argmax(y_test[i,:])
```

```
[13]: count=0
for i in range(10000):
    if(gt[i]==final_label[i]):
        count=count+1

print("Test set accuracy:",count/10000)
```

```
Test set accuracy: 0.5275
```



```
[0]: def shuffle_weights(model, weights=None):
    weights = model.get_weights()
    weights = [np.random.permutation(w.flat).reshape(w.shape) for w in weights]
    model.set_weights(weights)
```

### 3 Training the model with mean squared logarithmic loss

```
[21]: shuffle_weights(model)
model.compile(loss='mean_squared_logarithmic_error',
              optimizer=opt,
              metrics=['accuracy'])
history = model.fit(x_train,
    ↪y_train, batch_size=batch_size, epochs=20, validation_data=(x_test,
    ↪y_test), shuffle=True)
```

Train on 50000 samples, validate on 10000 samples

Epoch 1/20

50000/50000 [=====] - 63s 1ms/step - loss: 0.0047 - accuracy: 0.0619 - val\_loss: 0.0046 - val\_accuracy: 0.1335

Epoch 2/20

50000/50000 [=====] - 63s 1ms/step - loss: 0.0046 - accuracy: 0.1358 - val\_loss: 0.0044 - val\_accuracy: 0.1968

Epoch 3/20

50000/50000 [=====] - 63s 1ms/step - loss: 0.0044 - accuracy: 0.1914 - val\_loss: 0.0042 - val\_accuracy: 0.2454

Epoch 4/20

50000/50000 [=====] - 63s 1ms/step - loss: 0.0042 - accuracy: 0.2427 - val\_loss: 0.0040 - val\_accuracy: 0.2827

Epoch 5/20

50000/50000 [=====] - 63s 1ms/step - loss: 0.0041 - accuracy: 0.2796 - val\_loss: 0.0039 - val\_accuracy: 0.3196

Epoch 6/20

50000/50000 [=====] - 63s 1ms/step - loss: 0.0039 - accuracy: 0.3167 - val\_loss: 0.0038 - val\_accuracy: 0.3382

Epoch 7/20

50000/50000 [=====] - 63s 1ms/step - loss: 0.0038 - accuracy: 0.3433 - val\_loss: 0.0037 - val\_accuracy: 0.3756

Epoch 8/20

50000/50000 [=====] - 63s 1ms/step - loss: 0.0037 - accuracy: 0.3696 - val\_loss: 0.0036 - val\_accuracy: 0.3814

Epoch 9/20

50000/50000 [=====] - 63s 1ms/step - loss: 0.0036 - accuracy: 0.3952 - val\_loss: 0.0036 - val\_accuracy: 0.3949

Epoch 10/20

50000/50000 [=====] - 63s 1ms/step - loss: 0.0035 - accuracy: 0.4178 - val\_loss: 0.0035 - val\_accuracy: 0.4089

```

Epoch 11/20
50000/50000 [=====] - 63s 1ms/step - loss: 0.0034 -
accuracy: 0.4383 - val_loss: 0.0035 - val_accuracy: 0.4136
Epoch 12/20
50000/50000 [=====] - 63s 1ms/step - loss: 0.0033 -
accuracy: 0.4596 - val_loss: 0.0034 - val_accuracy: 0.4269
Epoch 13/20
50000/50000 [=====] - 63s 1ms/step - loss: 0.0032 -
accuracy: 0.4796 - val_loss: 0.0034 - val_accuracy: 0.4328
Epoch 14/20
50000/50000 [=====] - 63s 1ms/step - loss: 0.0031 -
accuracy: 0.4990 - val_loss: 0.0034 - val_accuracy: 0.4371
Epoch 15/20
50000/50000 [=====] - 62s 1ms/step - loss: 0.0030 -
accuracy: 0.5178 - val_loss: 0.0033 - val_accuracy: 0.4534
Epoch 16/20
50000/50000 [=====] - 63s 1ms/step - loss: 0.0029 -
accuracy: 0.5314 - val_loss: 0.0033 - val_accuracy: 0.4585
Epoch 17/20
50000/50000 [=====] - 63s 1ms/step - loss: 0.0028 -
accuracy: 0.5488 - val_loss: 0.0033 - val_accuracy: 0.4608
Epoch 18/20
50000/50000 [=====] - 63s 1ms/step - loss: 0.0027 -
accuracy: 0.5666 - val_loss: 0.0033 - val_accuracy: 0.4655
Epoch 19/20
50000/50000 [=====] - 63s 1ms/step - loss: 0.0026 -
accuracy: 0.5802 - val_loss: 0.0033 - val_accuracy: 0.4680
Epoch 20/20
50000/50000 [=====] - 63s 1ms/step - loss: 0.0026 -
accuracy: 0.5942 - val_loss: 0.0033 - val_accuracy: 0.4608

```

## 4 Plots for Mean Square Logarithmic Loss

Mean squared logarithmic error (MSLE) can be interpreted as a measure of the ratio between the true and predicted values. Mean squared logarithmic error is, as the name suggests, a variation of the Mean Squared Error.

```

[22]: # Generate generalization metrics
score = model.evaluate(x=x_test, y=y_test, verbose=1)
print(f'Test loss: {score[0]} / Test accuracy: {score[1]}')

# Visualize history
# Plot history: Loss
plt.plot(history.history['val_loss'])
plt.title('Validation loss history')
plt.ylabel('Loss value')
plt.xlabel('No. epoch')


```


```
plt.show()

# Plot history: Accuracy
plt.plot(history.history['val_accuracy'])
plt.title('Validation accuracy history')
plt.ylabel('Accuracy value (%)')
plt.xlabel('No. epoch')
plt.show()
```

10000/10000 [=====] - 5s 470us/step

Test loss: 0.003258741495758295 / Test accuracy: 0.4607999920845032





```
[0]: ans=model.predict(x_test)
final_label=np.zeros(10000)
gt=np.zeros(10000)

for i in range(10000):
    final_label[i]=np.argmax(ans[i,:])
    gt[i]=np.argmax(y_test[i,:])
```

```
[24]: count=0
for i in range(10000):
    if(gt[i]==final_label[i]):
        count=count+1

print("Test set accuracy:",count/10000)
```

Test set accuracy: 0.4608

```
[0]: def shuffle_weights(model, weights=None):
    weights = model.get_weights()
    weights = [np.random.permutation(w.flat).reshape(w.shape) for w in weights]
    model.set_weights(weights)
```

## 5 Using Hinge Loss for Training

```
[26]: shuffle_weights(model)
      model.compile(loss='hinge',
                    optimizer=opt,
                    metrics=['accuracy'])
      history = model.fit(x_train,
                          →y_train, batch_size=batch_size, epochs=20, validation_data=(x_test,
                          →y_test), shuffle=True)
```

Train on 50000 samples, validate on 10000 samples

Epoch 1/20

50000/50000 [=====] - 63s 1ms/step - loss: 1.0098 -  
accuracy: 0.0101 - val\_loss: 1.0098 - val\_accuracy: 0.0100

Epoch 2/20

50000/50000 [=====] - 62s 1ms/step - loss: 1.0098 -  
accuracy: 0.0099 - val\_loss: 1.0098 - val\_accuracy: 0.0100

Epoch 3/20

50000/50000 [=====] - 62s 1ms/step - loss: 1.0098 -  
accuracy: 0.0106 - val\_loss: 1.0097 - val\_accuracy: 0.0160

Epoch 4/20

50000/50000 [=====] - 63s 1ms/step - loss: 1.0095 -  
accuracy: 0.0235 - val\_loss: 1.0095 - val\_accuracy: 0.0240

Epoch 5/20

50000/50000 [=====] - 63s 1ms/step - loss: 1.0094 -  
accuracy: 0.0320 - val\_loss: 1.0093 - val\_accuracy: 0.0369

Epoch 6/20

50000/50000 [=====] - 63s 1ms/step - loss: 1.0093 -  
accuracy: 0.0361 - val\_loss: 1.0093 - val\_accuracy: 0.0341

Epoch 7/20

50000/50000 [=====] - 62s 1ms/step - loss: 1.0092 -  
accuracy: 0.0404 - val\_loss: 1.0090 - val\_accuracy: 0.0489

Epoch 8/20

50000/50000 [=====] - 62s 1ms/step - loss: 1.0091 -  
accuracy: 0.0453 - val\_loss: 1.0090 - val\_accuracy: 0.0497

Epoch 9/20

50000/50000 [=====] - 63s 1ms/step - loss: 1.0090 -  
accuracy: 0.0496 - val\_loss: 1.0089 - val\_accuracy: 0.0532

Epoch 10/20

50000/50000 [=====] - 63s 1ms/step - loss: 1.0090 -  
accuracy: 0.0519 - val\_loss: 1.0089 - val\_accuracy: 0.0560

Epoch 11/20

50000/50000 [=====] - 63s 1ms/step - loss: 1.0089 -  
accuracy: 0.0558 - val\_loss: 1.0088 - val\_accuracy: 0.0602

Epoch 12/20

50000/50000 [=====] - 62s 1ms/step - loss: 1.0088 -  
accuracy: 0.0619 - val\_loss: 1.0085 - val\_accuracy: 0.0724

```

Epoch 13/20
50000/50000 [=====] - 63s 1ms/step - loss: 1.0085 -
accuracy: 0.0731 - val_loss: 1.0084 - val_accuracy: 0.0785
Epoch 14/20
50000/50000 [=====] - 62s 1ms/step - loss: 1.0085 -
accuracy: 0.0765 - val_loss: 1.0085 - val_accuracy: 0.0775
Epoch 15/20
50000/50000 [=====] - 62s 1ms/step - loss: 1.0084 -
accuracy: 0.0809 - val_loss: 1.0082 - val_accuracy: 0.0908
Epoch 16/20
50000/50000 [=====] - 62s 1ms/step - loss: 1.0083 -
accuracy: 0.0832 - val_loss: 1.0082 - val_accuracy: 0.0900
Epoch 17/20
50000/50000 [=====] - 62s 1ms/step - loss: 1.0083 -
accuracy: 0.0862 - val_loss: 1.0081 - val_accuracy: 0.0969
Epoch 18/20
50000/50000 [=====] - 62s 1ms/step - loss: 1.0082 -
accuracy: 0.0891 - val_loss: 1.0083 - val_accuracy: 0.0871
Epoch 19/20
50000/50000 [=====] - 62s 1ms/step - loss: 1.0082 -
accuracy: 0.0883 - val_loss: 1.0081 - val_accuracy: 0.0933
Epoch 20/20
50000/50000 [=====] - 62s 1ms/step - loss: 1.0082 -
accuracy: 0.0903 - val_loss: 1.0082 - val_accuracy: 0.0897

```

## 6 Plots for Hinge Loss

Hinge loss is a loss function used for training classifiers. The hinge loss is used for “maximum-margin” classification, most notably for support vector machines (SVMs).

For an intended output  $t = \pm 1$  and a classifier score  $y$ , the hinge loss of the prediction  $y$  is defined as

$$\text{loss}(y) = \max(0, 1 - t \cdot y)$$

```

[27]: # Generate generalization metrics
score = model.evaluate(x=x_test, y=y_test, verbose=1)
print(f'Test loss: {score[0]} / Test accuracy: {score[1]}')

# Visualize history
# Plot history: Loss
plt.plot(history.history['val_loss'])
plt.title('Validation loss history')
plt.ylabel('Loss value')
plt.xlabel('No. epoch')
plt.show()


# Plot history: Accuracy


```

```
plt.plot(history.history['val_accuracy'])  
plt.title('Validation accuracy history')  
plt.ylabel('Accuracy value (%)')  
plt.xlabel('No. epoch')  
plt.show()
```

10000/10000 [=====] - 5s 468us/step

Test loss: 1.008201333618164 / Test accuracy: 0.08969999849796295





```
[0]: ans=model.predict(x_test)
final_label=np.zeros(10000)
gt=np.zeros(10000)

for i in range(10000):
    final_label[i]=np.argmax(ans[i,:])
    gt[i]=np.argmax(y_test[i,:])
```

```
[29]: count=0
for i in range(10000):
    if(gt[i]==final_label[i]):
        count=count+1

print("Test set accuracy:",count/10000)
```

Test set accuracy: 0.0897

```
[0]:
```